

## Deep learning for complete beginners: neural network fine-tuning techniques

Receive news and tutorials straight to your mailbox:

**SUBSCRIBE**

---

*[Edited on 20 March 2017, to account for API changes introduced by the release of Keras 2]*

---

### Introduction

Welcome to the third (and final) in a series of blog posts that is designed to get you quickly up to speed with *deep learning*; from first principles, all the way to

discussions of some of the intricate details, with the purposes of achieving respectable performance on two established machine learning benchmarks:

MNIST

(<http://yann.lecun.com/exdb/mnist>

(classification of handwritten digits) and CIFAR-10

(<https://www.cs.toronto.edu/~kriz/c>

(classification of small images across 10 distinct classes: plane, car, bird, cat, deer, dog, frog, horse, ship & truck).

# MNIST

# CIFAR-10



Last time around  
<https://cambridgespark.com/content/tutorials/neural-networks-tuning-techniques/index.html>,  
 I have introduced the *convolutional neural network* model, and illustrated how, combined with a simple but effective regularisation

method of *dropout*, it may quickly achieve an accuracy level of 78.6% on CIFAR-10, leveraging the Keras (<https://keras.io>) deep learning framework.

By now, you have acquired the fundamental skills necessary to apply deep learning to most problems of interest (a notable exception, outside of the scope of these tutorials, is the problem of processing *time-series of arbitrary length*, for which a *recurrent neural network* (RNN) model is often

preferable). In this tutorial, I will wrap up with an important but often overlooked aspect of tutorials such as this one – the tips and tricks for properly *fine-tuning* a model, to make it generalise better than the initial baseline you started out with.

This tutorial will, for the most part, assume familiarity with the previous two in the series.

---

Hyperparameter tuning and the baseline model

Typically, the design process for neural networks starts off by designing a simple network, either directly applying architectures that have shown successes ([http://rodrigob.github.io/are we t](http://rodrigob.github.io/are-we-t) for similar problems, or trying out hyperparameter values that generally seem effective. Eventually, we will hopefully attain performance values that seem like a nice baseline starting point, after which we may look into modifying every fixed detail in order to extract

the maximal performance capacity out of the network. This is commonly known as *hyperparameter tuning*, because it involves modifying the components of the network which need to be specified before training.

While the methods described here can yield far more tangible improvements on CIFAR-10, due to the relative difficulty of rapid prototyping on it without a GPU, we will focus specifically on

improving performance on the MNIST benchmark. Of course, I do invite you to have a go at applying methods like these to CIFAR-10 and see the kinds of gains you may achieve compared to the basic CNN approach, should your resources allow for it.

We will start off with the baseline CNN given below. If you find any aspects of this code unclear, I invite you to familiarise yourself with the previous two tutorials in



the series – all the relevant concepts have already been introduced there.

```
from keras.datasets
import mnist # subroutines
for fetching the MNIST
dataset

from keras.models import
Model # basic class for
specifying and training a
neural network

from keras.layers import
Input, Dense, Flatten,
Convolution2D,
MaxPooling2D, Dropout

from keras.utils import
np_utils # utilities for
one-hot encoding of ground
```

*truth values*

*batch\_size = 128 # in each iteration, we consider 128 training examples at once*

*num\_epochs = 12 # we iterate twelve times over the entire training set*

*kernel\_size = 3 # we will use 3x3 kernels throughout*

*pool\_size = 2 # we will use 2x2 pooling throughout*

*conv\_depth = 32 # use 32 kernels in both*

*convolutional layers*

*drop\_prob\_1 = 0.25 #*

*dropout after pooling with probability 0.25*

*drop\_prob\_2 = 0.5 # dropout in the FC layer with probability 0.5*

*hidden\_size = 128 # there will be 128 neurons in both hidden layers*

*num\_train = 60000 # there are 60000 training examples in MNIST*

*num\_test = 10000 # there are 10000 test examples in MNIST*

```
height, width, depth = 28,  
28, 1 # MNIST images are  
28x28 and greyscale  
num_classes = 10 # there  
are 10 classes (1 per  
digit)
```

```
(X_train, y_train),  
(X_test, y_test) =  
mnist.load_data() # fetch  
MNIST data
```

```
X_train =  
X_train.reshape(X_train.shape[0],  
height, width, depth)  
X_test =
```

```
X_test.reshape(X_test.shape[0],  
               height, width, depth)
```

```
X_train =
```

```
X_train.astype('float32')
```

```
X_test =
```

```
X_test.astype('float32')
```

```
X_train /= 255 # Normalise  
data to [0, 1] range
```

```
X_test /= 255 # Normalise  
data to [0, 1] range
```

```
Y_train =
```

```
np_utils.to_categorical(y_train,  
                        num_classes) # One-hot  
encode the labels
```

```
Y_test =
```

```
np_utils.to_categorical(y_  
    num_classes) # One-hot  
encode the labels
```

```
inp = Input(shape=(height,  
width, depth)) # N.B.  
TensorFlow back-end  
expects channel dimension  
last
```

```
# Conv [32] -> Conv [32] -  
> Pool (with dropout on  
the pooling layer)
```

```
conv_1 =  
Convolution2D(conv_depth,  
(kernel_size,  
kernel_size),
```